OpenFOS

Open Founder Operating System

Master Skill Map v2.0

24 Skills + 2 Reference Modules
Organized by FOUNDRY Framework

Seranai · March 2026

Created by Arshad Sayyad · A Seranai Product · Licensed CC BY-NC-SA 4.0

What Is This Skill?

Experiment Design & Rapid Testing is an AI-powered tool that helps you answer the most important question a founder can ask after validating a problem: How do I prove there is real demand before I invest months building?

It works as a conversation. You describe your idea or hypothesis in plain language and the skill guides you through a structured experiment design process used by the best lean startup operators in the world. No templates to fill out, no jargon to learn, no code required.

THE CORE STATISTIC
Founders who test before they build are 3x more likely to find product-market fit.
This skill exists to make sure you are testing the right things, in the right order, with the right success criteria.

By the end of one conversation (about 15 minutes), you will have a ranked assumption map, a complete experiment design brief, and a concrete kill condition that tells you exactly when to stop and pivot.

Who Is This For?

This skill is designed for founders who have already identified a problem worth solving and now need to validate demand before committing engineering time and capital. It is especially useful if:

- You have a hypothesis and want to design the cheapest test that produces real signal
- You have already run one experiment and need to interpret results and decide what comes next
- You're tempted to start building but know you should validate first
- You're choosing between two possible experiments and want a structured way to pick
- You're a first-time founder who has never run a structured lean experiment before
- You're in an emerging market and need experiment formats that work without high digital ad spend or card penetration

It works for founders anywhere in the world. Whether you're testing in San Francisco, Lagos, Bangalore, Dubai, or São Paulo, the experiment framework adapts to your channel access, budget, and customer context.

How to Get Started

Using this skill requires three things: access to Claude (claude.ai or the Claude app), the Experiment Design skill installed, and a hypothesis or idea described in your own words.

Step 1: Open Claude and Start a New Conversation

Go to claude.ai or open the Claude app. Start a new conversation. That's it.

Step 2: Describe Your Hypothesis or Idea

Type your startup hypothesis, experiment question, or idea in plain language. You do not need a formatted hypothesis or a prior experiment. Just describe what you want to test and who you think will care about it.

EXAMPLE 1 — Detailed Hypothesis
"I want to test whether VP-level operations leaders at logistics companies will pay $400/month for a tool that automatically flags supplier contract compliance issues. I believe they spend 8+ hours a week doing this manually today."

EXAMPLE 2 — Brief
"I think HR teams at mid-size companies would pay to automate onboarding document collection. I want to run a quick test."

EXAMPLE 3 — Minimal
"Invoice automation for freelance designers."

All three will work. The more detail you provide, the sharper the initial analysis. But even a single sentence is enough to get started — the skill will ask targeted questions to fill in the gaps.

Step 3: Follow the Conversation

The skill will guide you through four phases. You do not need to remember them or do anything special — just respond naturally to what it asks. The whole process takes about 15 minutes.

What Happens: The Four Phases

The skill walks you through four phases, each building on the last. Here is what to expect at each stage.

PHASE 1: ASSUMPTION MAP + EXPERIMENT SHORTLIST *Your idea gets an immediate intelligence briefing • ~30 seconds*	**What You Do** • Describe your hypothesis or idea (the skill does the rest)	**What You Get** • Ranked list of your top 3 untested assumptions • Experiment Readiness Score (Clarity + Evidence = 1-10) • Top 3 experiment recommendations ordered by cost-to-signal ratio • Time, cost, and success criteria for each experiment

WHY THIS MATTERS:

Most founders jump straight to building a landing page. This skill tells you whether a landing page is even the right test — or whether a 20-minute phone call would give you better signal in less time for zero dollars.

PHASE 2: SHARPENING QUESTIONS *Filling in the gaps to sharpen the experiment design • ~3 minutes*	**What You Do** • Answer 3-4 targeted questions about your target customer, access channels, budget, and success criteria • It's okay to say "I don't know" — that's valuable data	**What You Get** • A much sharper understanding of what experiment to run and what success looks like

The questions are never generic. They are selected based on what is missing from your initial description. If you already described your target customer in detail, the skill will not ask you about your customer again. It will focus on what is missing.

PHASE 3: FULL EXPERIMENT DESIGN BRIEF *The complete, ready-to-run experiment • ~5 minutes to read*	**What You Do** • Answer the sharpening questions from Phase 2	**What You Get** • Formal hypothesis statement (We believe / If we / Then) • Step-by-step experiment execution plan • Exact success threshold and kill condition • Tools, scripts, and outreach templates to run it immediately

Created by Arshad Sayyad · A Seranai Product · Licensed CC BY-NC-SA 4.0

PHASE 4: RESULTS INTERPRETATION *Turning data into your next decision • ~3 minutes to read*	**What You Do** • Share your experiment results (numbers, responses, or outcomes)	**What You Get** • Signal quality rating on the 5-level evidence ladder • GO / ITERATE / PIVOT / KILL recommendation with specific reasoning • If iterating: what to change in the next experiment • If GO: a bridge to the next step (Persona Definition or Product Development)

Understanding Your Results

The Experiment Readiness Score

This score tells you how much signal your experiment is likely to produce before you run it. It is calculated across two dimensions:

Dimension	What It Measures
Clarity Score (1-5)	How precisely defined is your target customer and problem? (1 = "anyone with a problem" / 5 = a specific job title, company size, and frequency of pain)
Evidence Score (1-5)	How much existing signal already exists that the problem is real? (1 = founder intuition only / 5 = paying customers plus confirmed interview data)
Readiness Score	Clarity + Evidence = 1-10. Below 5 means your experiment results will likely be noisy. The skill flags this and recommends sharpening first.

The Signal Quality Ladder

Not all evidence is equal. After your experiment, the skill rates the quality of your results on this five-level ladder:

Level	Signal Type	Example	Reliability
1 - Noise	Verbal enthusiasm from friends/family	*"That's a great idea!"*	Very Low
2 - Weak	Social engagement from strangers	*Likes, shares, retweets on a post*	Low
3 - Mo	Time investment	*Completing a survey or attending a demo*	Moderate

derate	from target customers		
4 - Strong	Data or access shared	*Providing API keys, internal data, referrals*	High
5 - Validated	Money exchanged	*Pre-order, deposit, LOI signed, invoice paid*	Very High

The Verdict: GO / ITERATE / PIVOT / KILL

🟢 GO — Strong signal. Keep building.

Your experiment passed its success threshold. You have Level 3 signal with N>15, Level 4 signal with N>5, or at least one paying stranger. Proceed with confidence to the next phase: customer interviews for qualification, persona definition, or product development.

🟡 ITERATE — Signal exists but the design needs refinement.

Your experiment showed real interest but did not cross the success threshold. The skill will tell you exactly what to change in the next experiment: a different value proposition, a sharper target segment, a different channel, or a different test type. Run the revised experiment before drawing conclusions.

🟠 PIVOT — The core insight has merit, but the framing is wrong.

The riskiest assumption failed, but there is evidence pointing to a different customer, problem angle, or price point worth testing. The skill will specify the pivot: do not abandon the domain, but reframe and run a new experiment against the updated hypothesis.

🔴 KILL — The evidence does not support this hypothesis.

Multiple assumptions failed and no pivot direction has clear signal. This is not failure — killing a bad hypothesis early is one of the most valuable things a founder can do. You just saved months of wasted build time. The skill will suggest where to redirect your energy based on what it learned about your strengths and market observations.

The Experiment Menu

The skill matches experiments to your stage, resources, and channel access. Here is a summary of the core experiments and when to use each.

Experiment	Best For	Cost	Minimum Signal to Proceed
Customer Discovery Interviews	Does this problem exist and how painful is it?	Free	*15 interviews; 3+ unprompted referrals or purchase offers*
Concierge MVP	Will they pay for the value before you build it?	Free (your time)	*Customers pay AND re-engage*
LinkedIn / Community Post	Does the problem language resonate with strangers?	Free	*>20 meaningful DMs or replies from target customers*
Fake Door Landing Page	Is there demand without building anything?	$0-$50/yr	*>20% visitor-to-signup conversion*
Clickable Prototype Test	Does the solution approach make intuitive sense?	$0-$99/mo	*Users navigate + ask "when can I use this?"*
Smoke Test / Paid Ad Campaign	Does the value proposition resonate with strangers?	$100-$300	*CTR >2% on LinkedIn; LP conversion >20%*
Wizard of Oz Test	Can you simulate a complex product before it exists?	$50-$200	*Retention and payment after first "automated" experience*
Pre-Sale / Deposit Test	Will anyone pay money before the product exists?	Free to set up	*Any paying stranger; 5+ to proceed with high confidence*

Tips for Getting the Most Out of This Skill

Test the riskiest assumption first, not the easiest one. The cheapest experiment is not always the best one. The best experiment is the one that tests your riskiest assumption. If you have no idea whether customers want this at all, do not start with a paid ad campaign. Start with 15 phone calls.

"I don't know" is a valid answer. If you have not talked to customers yet, say so. If you are not sure about pricing, say so. The skill treats gaps in knowledge as data points — they tell you exactly what to test.

Charge something. Even $1. Free users give false signal. The moment money changes hands, you are measuring revealed preference — what people actually value — rather than stated preference — what they say they might want.

Do not show your experiment only to friends and family. Their enthusiasm is noise. Their conversion rate is meaningless. You need signal from strangers in your target segment.

Run it in 7 days, not 30. A well-designed experiment should produce usable signal within one week. If your experiment requires 30 days to get results, it is too long. Shorten it or change the test type.

Challenge a KILL signal with evidence, not conviction. If the skill recommends KILL and you disagree, that is okay. Founder conviction matters. But use the revised experiment plan to gather the specific evidence that would change the score. Do not just disagree — prove it.

What Comes Next

Experiment Design & Rapid Testing is Skill 12 of 12 in the OpenFOS sequence. Once you have strong experiment signal (GO or successful ITERATE), the natural next step depends on where you are in your journey:

#	Skill	Question	Status
1	Problem Validation	Is this worth building?	Completed
2	TAM-SAM-SOM	How big is the opportunity?	Completed
3	Market & Beachhead	Where do I compete first?	Completed
4	Persona & Testing	Who exactly buys this?	Completed
12	**Experiment Design & Rapid Testing**	How do I prove demand fast?	✅ You are here

To move to the next phase, simply start a new conversation and say: "I ran an experiment on [your hypothesis] and here are the results." The skill will pick up from there and tell you exactly what to do next.

Frequently Asked Questions

Do I need any technical skills to use this?

No. You describe your hypothesis in plain language and have a conversation. The skill generates the experiment design, tools list, and outreach templates for you.

How long does it take?

About 15 minutes for the full conversation, depending on how much detail you provide in your first message.

What if I have already run an experiment and just want help interpreting results?

Perfect. Just share your results in your first message and the skill will skip directly to Phase 4: Results Interpretation. You do not have to go through the full design process again.

Can I use this for B2C ideas?

The skill is optimized for B2B, where buying dynamics, pricing models, and validation methods differ significantly from B2C. A B2C-adapted version is on the roadmap.

What if I disagree with a KILL recommendation?

Good. That means you have conviction. Use the revised experiment plan the skill provides to gather specific evidence that would change the score. If you cannot find that evidence, the score was probably right.

I'm in an emerging market. Will this work for me?

Yes. The skill includes specific adjustments for founders building in or for emerging markets, including lower digital ad efficiency, lower card penetration for pre-sales, and WhatsApp or in-person experiment formats as alternatives.

Is my data private?

Your conversation is between you and Claude. The skill does not store or share your idea with anyone.

The FOUNDRY Framework

FOUNDRY is the sequential journey framework that guides founders from initial idea through scale or exit. Each letter represents a stage, and each stage contains specific skills a founder needs to master.

Stage	Purpose	Skills
F – Foundation	Validate the idea. Confirm the problem is real, the market exists, and you know who to serve.	Skills 1–4
O – Outline	Design the business. Articulate value, map competition, model revenue, scope the MVP.	Skills 5–8
U – User	Acquire and understand customers. Discover needs, design GTM, set pricing, run experiments.	Skills 9–12
N – Nurture	Build relationships and traction. Craft the pitch, model financials, build the founder narrative.	Skills 13–15
D – Discipline	Execute and operate. Set OKRs, hire the team, handle legal and entity setup.	Skills 16–18
R – Results	Measure and iterate. Track KPIs, assess PMF, decide whether to pivot or persevere.	Skills 19–20
Y – Yield	Scale or exit. Confirm scaling readiness, plan liquidity events.	Skills 21–22

Cross-cutting skills (23–24) are not stage-specific and can be invoked at any point in the journey.

Reference modules (R1–R2) are shared evaluation lenses that multiple skills invoke. They are never used in isolation.

Created by Arshad Sayyad · A Seranai Product · Licensed CC BY-NC-SA 4.0